

# Homework 7

due Wed Nov 2 at start of class

Homework7.pdf

HW7.ipynb or HW7.py

download from Brightspace

Images.zip

# **Images and Signals**

and some basics of image and signal processing

# **PIL**

Pillow - an updated fork of PIL (Python Imaging Library)

<https://pillow.readthedocs.io/en/stable/>

# opening with PIL

```
from PIL import Image
```

```
im = Image.open( 'Jordinggray.bmp' )
```



**object with image attributes and methods**

```
im.show( )
```

**opens image in a new window**

```
from IPython.display import display
```

```
display(im)
```

**opens image in Jupyter Notebook**

# manipulating images with PIL

```
im.thumbnail((100,100))
```

**create a thumbnail**  
**(manipulates image itself)**

```
im2 = im.rotate(45)
```

**rotates images**  
**(returns new image)**

# more manipulating images with PIL

```
from PIL import ImageEnhance
```

```
im2 = ImageEnhance.Contrast(im)  
display(im2.enhance(4))
```

```
im2 = ImageEnhance.Brightness(im)  
display(im2.enhance(.4))
```

```
im2 = ImageEnhance.Sharpness(im)  
display(im2.enhance(2.0))
```

# saving images with PIL

```
im = Image.open( 'Jordinggray.bmp' )
```

```
im2 = im.rotate(45, expand=True)
```

```
im2.save( 'Jordinggray_Rot45.jpg', format='JPEG' )
```



# extracting numpy data from images

```
im = Image.open( 'Jordingray.bmp' )
```

```
imdata = np.asarray(im)
```

```
print( type(imdata) )
```

<class 'numpy.ndarray'>

```
print( type(imdata[0,0]) )
```

<class 'numpy.uint8'>

```
print( imdata.shape )
```

(512, 512)

create an image from a numpy array

```
im2 = Image.fromarray(imdata)
```

```
im2.save('Jordinggray_New.jpg', format='JPEG')
```

# manipulating image data in numpy arrays

unsigned int arithmetic

255 = 11111111  
decimal          binary

$$2^7 + 2^6 + 2^5 + 2^4 + 2^3 + 2^2 + 2^1 + 2^0 = 255$$
$$128 + 64 + 32 + 16 + 8 + 4 + 2 + 1 = 255$$

# manipulating image data in numpy arrays

$$\begin{array}{ccccccc} 255 & + & 1 & = & ? \\ \text{uint8} & & \text{uint8} & & \text{uint8} \end{array}$$

# manipulating image data in numpy arrays

255     +     1     =     0     Python  
uint8     uint8     uint8

## Matlab

```
Command Window
>> uint8(255) + uint8(1)
ans =
    uint8
    255
fx >>
```

# manipulating image data in numpy arrays

255 + 2 = 1 Python  
uint8 uint8 uint8

Matlab

```
Command Window
>> uint8(255)+uint8(2)

ans =

    uint8
    255

fx >>
```

# manipulating image data in numpy arrays

255    +    100    =    99    Python  
uint8        uint8        uint8

## Matlab

Command Window

```
>> uint8(255)+uint8(100)
```

```
ans =
```

```
uint8
```

```
255
```

```
fx >> |
```

# manipulating image data in numpy arrays

```
imdata = np.asarray(im)
```

```
imdata2 = imdata + 150
```





# manipulating image data in numpy arrays

```
imdata = np.asarray(im)
```

```
imdata2 = imdata + 255
```



# manipulating image data in numpy arrays

```
imdata = np.asarray(im)
```

```
imdata2 = imdata + 1000
```



?

# manipulating image data in numpy arrays

```
imdata = np.asarray(im)
```

```
      uint16      uint8      uint16  
imdata2 = imdata + 1000
```

```
imdata2
```

```
array([[1091, 1092, 1093, ..., 1106, 1109, 1109],  
      [1084, 1085, 1087, ..., 1110, 1110, 1110],  
      [1085, 1083, 1082, ..., 1108, 1108, 1109],  
      ...,  
      [1169, 1175, 1179, ..., 1184, 1184, 1187],  
      [1171, 1177, 1178, ..., 1184, 1185, 1185],  
      [1173, 1177, 1178, ..., 1184, 1184, 1184]] , dtype=uint16)
```

# manipulating image data in numpy arrays

```
imdata = np.asarray(im)
```

```
      uint16      uint8      uint16  
imdata2 = imdata + 1000
```

```
imdata2
```

```
array([[1091, 1092, 1093, ..., 1106, 1109, 1109],  
      [1084, 1085, 1087, ..., 1110, 1110, 1110],  
      [1085, 1083, 1082, ..., 1108, 1108, 1109],  
      ...,  
      [1169, 1175, 1179, ..., 1184, 1184, 1187],  
      [1171, 1177, 1178, ..., 1184, 1185, 1185],  
      [1173, 1177, 1178, ..., 1184, 1184, 1184]], dtype=uint16)
```

```
im2 = Image.fromarray(imdata2)
```

```
im2.save('tmp.jpg', format='JPEG') Error
```

## converting to/from floats

```
np.uint8(np.float64(np.uint8(255)+np.uint8(100)))
```

```
[140]: 99
```

# converting to/from floats

```
np.uint8(np.float64(np.uint8(255)+np.uint8(100)))
```

```
[140]: 99
```

```
imdataf = np.float64(imdata)
```

```
imdataf = imdataf*500 + 250
```


```
imdata2 = np.uint8((imdataf/np.max(imdataf))*255)
```



# image histograms


## greyscale image

```
hist = np.array(im.histogram())  
pix = np.arange(256)  
plt.bar(pix, hist, width=1, alpha=.3, color='grey')
```



## color image

```
hist = np.array(im.histogram())  
hist = np.resize(hist, (3,256))  
pix = np.arange(256)  
plt.bar(pix, hist[0,:], width=1, alpha=.3, color='red')  
plt.bar(pix, hist[1,:], width=1, alpha=.3, color='green')  
plt.bar(pix, hist[2,:], width=1, alpha=.3, color='blue')
```



# simple image processing

a very basic operation:

convert an image to a pure black (0) or white (255) image, using a criterion as the cut off

how would we do this?



# simple image processing

```
im = Image.open( 'Jordingray.bmp' )
```

```
imdata = np.asarray(im)
```

```
imdata2 = np.uint8(255*(imdata>C))
```

example of "vectorized" code

```
im2 = Image.fromarray(imdata2)
```

```
display(im2)
```

# simple image processing

that's a step function ...

what about a more general function?

$$\psi(x) = \gamma + (1 - \gamma - \lambda) F(x)$$

$$F(x) = 1 - \exp\left[-\left(\frac{x}{\alpha}\right)^\beta\right]$$

should look  
familiar

how could we modify this?


# simple image processing

$\alpha=100$

$\beta=3$

$\gamma=0.0$

$\lambda=0.0$




$\alpha=100$

$\beta=50$

$\gamma=0.0$

$\lambda=0.0$




# simple image processing

$\alpha=100$

$\beta=3$

$\gamma=0.0$

$\lambda=0.0$




$\alpha=100$

$\beta=50$

$\gamma=0.0$

$\lambda=0.0$



0.0-1.0

0-255

# simple image processing

```
imdata2 = myfilter(imdata, alpha=100, beta=3,  
                   gam=0.0, lam=0.0)
```

**need to make sure imdata2 is `uint8` between 0 and 255**

# simple image processing

```
imdata2 = myfilter(imdata, alpha=100, beta=3,  
                  gam=0.0, lam=0.0)
```

**need to make sure imdata2 is uint8 between 0 and 255**

```
imdata2 = np.uint8(255*imdata2)
```

```
im2 = Image.fromarray(imdata2)
```

**original**



**filtered**



# simple image processing

```
imdata2 = myfilter(imdata, alpha=150, beta=8,  
                  gam=0.4, lam=0.4)
```

need to make sure imdata2 is `uint8` between 0 and 255

```
imdata2 = np.uint8(255*imdata2)
```

```
im2 = Image.fromarray(imdata2)
```

original




filtered



a different transformation applied to different pixels

# image masks



how could we carry out this operation?




# image masks

create mask

```
sz = 250
```

```
mask = np.zeros((r,c))
```

```
mask[int(r/2-sz/2):int(r/2+sz/2),  
      int(r/2-sz/2):int(r/2+sz/2)] = 1
```



# image masks

create mask

```
sz = 250
```


```
mask = np.zeros((r,c))
```

```
mask[int(r/2-sz/2):int(r/2+sz/2),  
      int(r/2-sz/2):int(r/2+sz/2)] = 1
```

```
imdata2 = np.uint8(imdata * mask)
```



**x**




**=**




# image masks



**x**



**=**



how could we carry out this operation?



# image masks

## create mask

```
def my_mvnpdf(x, M, S):  
    return (1 / np.sqrt((2 * np.pi * np.linalg.det(S)))) *  
    np.exp(-(1/2) * (np.transpose(x-M) @ np.linalg.inv(S) @  
    (x-M)))
```

```
mx = r/2; my = c/2  
sx = 50; sy = 50; rho = 0.0  
mu = np.array([mx, my])  
sig = np.array([[sx**2,      rho*sx*sy],  
                [rho*sx*sy, sy**2]])
```

**basically, same code from before (with 3D plots)**



# image masks

create mask

```
inc = 1
xsteps = np.arange(0, r, 1)
ysteps = np.arange(0, c, 1)
X, Y = np.meshgrid(xsteps, ysteps)
Z = np.zeros(X.shape)

for i in range(len(xsteps)):
    for j in range(len(ysteps)):
        xy = np.array([X[i,j], Y[i,j]])
        Z[i,j] = my_mvnpdf(xy, mu, sig)
```

```
mask = Z/np.max(Z)
```

*no longer treating this as a probability distribution,  
we want the center point intensity equal to 1.0,  
do not care that the area under the curve equals 1.0*

# image masks

apply mask

```
imdata2 = np.uint8(imdata * mask)

im2 = Image.fromarray(np.uint8(imdata2))
display(im2)
```

# image processing scripts

```
import os
import re

pattern = r'.*((jpg)|(JPG))'
dir = './images'
with os.scandir(dir) as entries:
    for entry in entries:
        if re.search(pattern, entry.name):
            fullpath = dir + '/' + entry.name
            im = Image.open(fullpath)
            im.thumbnail((100,100))
            display(im)
```

of course, this could do  
something more complicated  
than display the images





# **More on Images and Signals**

our data structures are discrete in time or space

time = recorded sounds (world), recorded voltages (brain)

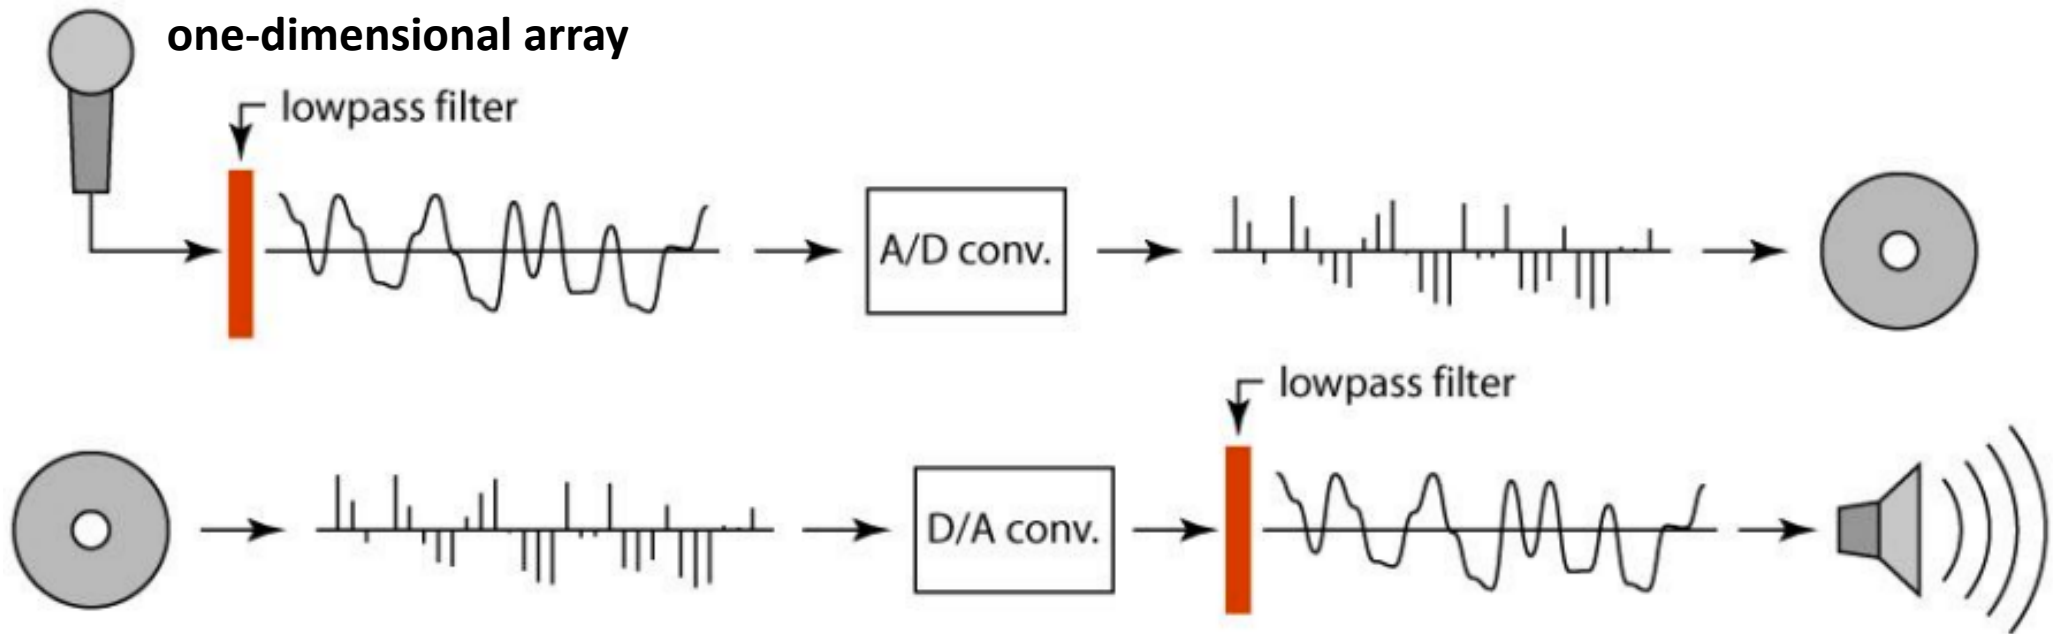
space = images (objects, scenes, brains)

in the real world, time and space are continuous, not discrete

in mathematics, there are continuous and discrete versions


**Audio**

**one-dimensional array**



**B/W image**

**two-dimensional array**



|        |        |        |        |        |        |        |
|--------|--------|--------|--------|--------|--------|--------|
| 0.2251 | 0.2563 | 0.2826 | 0.2826 | 0.4    |        |        |
| 0.5342 | 0.2051 | 0.2157 | 0.2826 | 0.3822 | 0.4391 | 0.4391 |
| 0.5342 | 0.1789 | 0.1307 | 0.1789 | 0.2051 | 0.3256 | 0.2483 |
| 0.4308 | 0.2483 | 0.2624 | 0.3344 | 0.3344 | 0.2624 | 0.2549 |
| 0.3344 | 0.2624 | 0.3344 | 0.3344 | 0.3344 | 0.3344 | 0.3344 |

**color image**

**three-dimensional array**

**video**

**four-dimensional array**

**video + sound**

**five-dimensional array**

# key factors that impact quality of digitization

**spatial resolution / sampling rate:** how close together in space (image/video) or time (audio) samples are taken

**bit depth:** what range of values (intensities) can be registered depending on the number of bits/bytes devoted to the signal

**spatial resolution / sampling rate**


# spatial resolution / sampling rate

CD Audio: 44.1kHz\* (44,100 samples per second)

HD Audio: 96kHz (96,000 samples per second)


Blu Ray Audio: up to 192kHz

\* minimal sampling rate to reproduce 20kHz sound frequencies  
(general limit of human hearing)



if you were recording or playing sounds for non-human animals  
you would want to make sure you had a sufficient sampling rate


# spatial resolution / sampling rate



8K is 4x  
as large



# spatial resolution / sampling rate



**bit depth**

# bit depth

double-precision floating point



single-precision floating point



8-bit signed integer



16-bit signed integer



16-bit unsigned integer



64-bit unsigned integer



# bit depth (audio)

CD Audio: 16 bits (2 bytes) per sample

DVD/Blu Ray: 24 (3 bytes) per sample

Signal-to-noise ratio and resolution of bit depths

| # bits | SNR       | Possible integer values (per sample) | Base-ten signed range (per sample)                       |
|--------|-----------|--------------------------------------|--|
| 4      | 24.08 dB  | 16                                   | −8 to +7   |
| 8      | 48.16 dB  | 256                                  | −128 to +127   |
| 11     | 66.22 dB  | 2048                                 | −1024 to +1023   |
| 12     | 72.24 dB  | 4096                                 | −2048 to +2047   |
| 16     | 96.33 dB  | 65,536                               | −32,768 to +32,767                                       |
| 20     | 120.41 dB | 1,048,576                            | −524,288 to +524,287                                     |
| 24     | 144.49 dB | 16,777,216                           | −8,388,608 to +8,388,607                                 |
| 32     | 192.66 dB | 4,294,967,296                        | −2,147,483,648 to +2,147,483,647                         |
| 48     | 288.99 dB | 281,474,976,710,656                  | −140,737,488,355,328 to +140,737,488,355,327             |
| 64     | 385.32 dB | 18,446,744,073,709,551,616           | −9,223,372,036,854,775,808 to +9,223,372,036,854,775,807 |

## **bit depth (images)**

|                |  |
|----------------|--|
| gif, jpg, tiff | 8 bits per channel (R, G, B) - 24 bits total |
| raw            | 12 or 14 bits per channel - 36-42 bits total |

# bit depth (monitors)

1st computers

1 bit (white/black, green/black)

early NeXT

2 bits

Atari, Commodore 64

4 bits

used color maps

early VGA

8 bits (total for color)

used color maps

most current monitor

24 bits (R, G, B)

HDR monitors

30/36/48 bits (R, G, B)

more than 3 primaries

R, G, B, C, Y, M